

Software Engineering

Cicio Ionuț

22/12/2024

Contents

1 Software models	4
1.1 The “Amazon Prime Video” article	4
1.2 Formal notation	5
1.2.1 The concept of time	5
1.2.2 Markov Chain	5
1.2.3 DTMC (Discrete Time Markov Chain)	5
1.2.3.1 An example of DTMC	7
1.2.4 Network of Markov Chains	7
1.3 Tips and tricks	7
1.3.1 Mean	7
1.3.2 Eulero’s method for differential equations	8
2 C++	9
2.1 Intro to <code>#include <random></code>	9
2.1.1 Seed & <code>std::default_random_engine</code>	9
2.1.2 Distributions	9
2.1.2.1 <code>std::uniform_int_distribution<>()</code>	9
2.1.2.2 <code>std::uniform_real_distribution<>()</code>	9
2.1.2.3 <code>std::bernoulli_distribution<>()</code>	9
2.1.2.4 <code>std::poisson_distribution<>()</code>	9
2.1.2.5 <code>std::geometric_distribution<>()</code>	9
2.1.2.6 <code>std::discrete_distribution<>()</code>	9
2.2 Intro to data structures	9
2.2.1 <code>std::vector<T>()</code>	9
2.2.2 <code>std::deque<T>()</code>	9
2.2.3 Sets	9
2.2.4 Maps	9
2.3 I/O	9
2.3.1 <code>#include <iostream></code>	9
2.3.2 Files	9
3 Exercises	10
3.1 First examples (1000)	10
3.1.1 A simple Markov Chain (1100)	10
3.1.2 Connect Markov Chains pt.1 (1200)	11
3.1.3 Connect Markov Chains pt.2 (1300)	11
3.1.4 Connect Markov Chains pt.3 (1400)	11
3.2 Traffic light (2100, 2200, 2300)	11
3.3 Control center (3000)	12
3.3.1 No network (3100)	12
3.3.2 Network monitor (no faults) (3200)	12
3.3.3 Network monitor (faults, no repair) (3300)	12
3.3.4 Network monitor (faults, repair) (3400)	12

3.3.5 Network monitor (faults, repair, correct protocol) (3500)	12
3.4 Statistics (4000)	12
3.4.1 Expected value (4100)	12
3.4.2 Probability (4200)	12
3.5 Transition matrix	12
3.6 Complex systems	12
3.6.1 Insulin pump	12
3.6.2 Buffer	12
3.6.3 Server	12
4 Exam	12
4.1 Development team (time & cost)	12
4.2 Backend load balancing	12
4.3 Heater simulation	12
Bibliography	13

1 Software models

When designing **complex software** we have to make major **design choices** at the beginning of the project. Often those choices can't be driven by experience or reasoning alone, that's why a **model** of the project is needed to compare different solutions. Our formal tool of choice is the **Discrete Time Markov Chain** (Section 1.2.3).

1.1 The “Amazon Prime Video” article

If you were tasked with designing the software architecture for **Amazon Prime Video** (*a live streaming service for Amazon*), how would you go about it? What if you had the **non-functional requirement** to keep the costs as low as possible?

In a recent article, Marcin Kolny, a Senior SDE at Prime Video, describes how they “*reduced the cost of the audio/video monitoring infrastructure by 90%*” [1] by using a monolith application instead of distributed microservices (an outcome one wouldn't usually expect).

While there isn't always definitive answer, one way to go about this kind choice is building a model of the system to compare the solutions. In the case of Prime Video, “*the audio/video monitoring service consists of three major components:*” [1]

- the **media converter** converts input audio/video streams
- the **defect detectors** execute algorithms that analyze frames and audio buffers in real-time looking for defects and send notifications
- the **orchestrator** controls the flow in the service

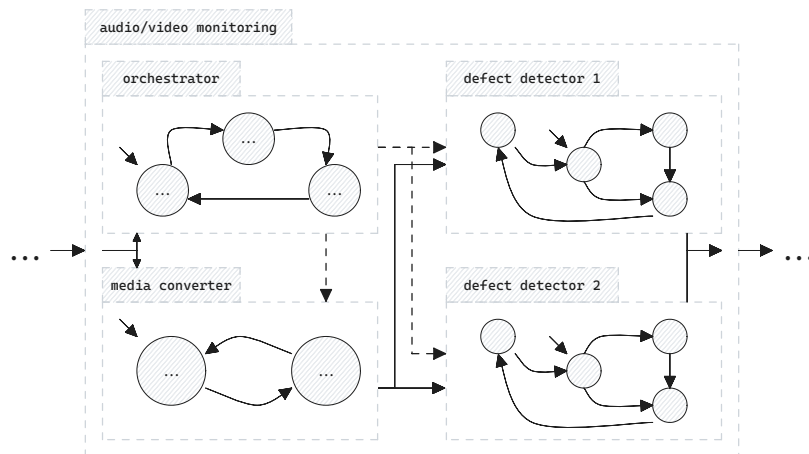


Figure 1: Model of the audio/video monitoring system

We want to model the components as some kind of stateful machine (with inputs and outputs, Figure 1), interconnect them and **simulate** the behaviour of the system as a whole.

1.2 Formal notation

1.2.1 The concept of time

TODO: rewrite

The models treated in the course evolve through **time**. Time can be modeled in many ways (I guess?), but, for the sake of simplicity, we will consider discrete time. Let W be the '*weather system*' and D the '*driving ability*' system in Section 1.2, we can define the evolution of D as

$$\begin{aligned} D(0) &= \text{'good'} \\ D(t+d) &= f(D(t), W(t)) \end{aligned} \tag{1}$$

Given a time instant t (let's suppose 12:32) and a time interval d (1 minute), the driving ability of D at 12:33 depends on the driving ability of D at the time 12:32 and the weather at 12:32.

1.2.2 Markov Chain

A Markov Chain is...

1.2.3 DTMC (Discrete Time Markov Chain)

A DTMC M is a tuple (U, X, Y, p, g) s.t.

- $U \neq \emptyset \wedge X \neq \emptyset \wedge Y \neq \emptyset$ (*otherwise stuff doesn't work*)
- U can be either
 - $\{u_1, \dots, u_n\}$ where u_i is an input value
 - $\{()\}$ if M doesn't take any input
- $X = \{x_1, \dots, x_n\}$ where x_i is a state
- $Y = \{y_1, \dots, y_n\}$ where y_i is an output value
- $p : X \times X \times U \rightarrow [0, 1]$ is the transition function
- $g : X \rightarrow Y$ is the output function

$$\forall x \in X \quad \forall u \in U \quad \sum_{x' \in X} p(x'|x, u) = 1 \tag{2}$$

$$\begin{aligned} M(0) &= x_1 \\ M(t+d) &= \begin{cases} x_1 & \text{with probability } p(x_1|M(t), U(t)) \\ x_2 & \text{with probability } p(x_2|M(t), U(t)) \\ \dots & \end{cases} \end{aligned} \tag{3}$$

It's interesting to notice that the transition function depends on the input values. If you consider the '*driving ability*' system in Section 1.2, you can see that the probability to go from **good** to **bad** is higher if the weather is rainy and lower if it's sunny.

1.2.3.1 An example of DTMC

Let's consider the development process of a team. We can define a DTMC $M = (U, X, Y, p, g)$ s.t.

- $U = \{()\}$, as it doesn't have any input
- $X = \{0, 1, 2, 3\}$
- $Y = \text{Cost} \times \text{Duration (in months)}$

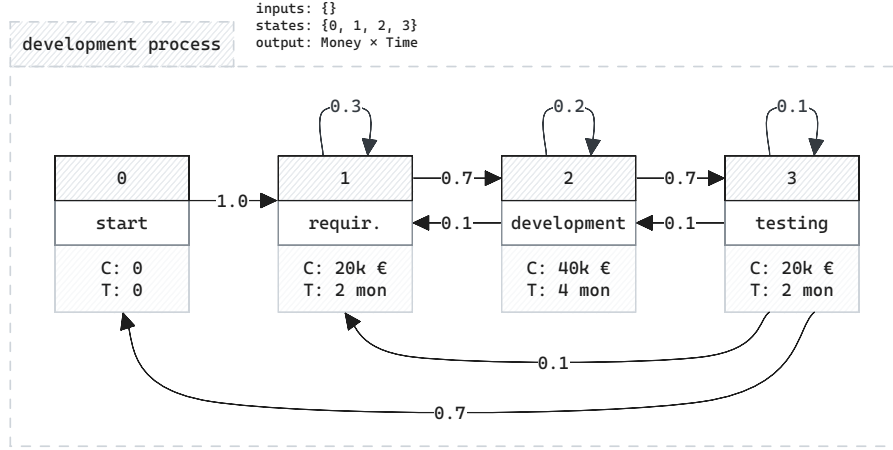


Figure 2: the model of a team's development process

$$g(x) = \begin{cases} (0, 0) & \text{if } x = 0 \\ (20000, 2) & \text{if } x = 1 \\ (40000, 4) & \text{if } x = 2 \\ (20000, 2) & \text{if } x = 3 \end{cases} \quad (4)$$

1.2.4 Network of Markov Chains

TODO...

1.3 Tips and tricks

1.3.1 Mean

TODO: 'mean' trick, ggwp

$$\begin{aligned} \varepsilon_n &= \frac{\sum_{i=0}^n v_i}{n} \\ \varepsilon_{n+1} &= \frac{\sum_{i=0}^{n+1} v_i}{n+1} = \frac{\left(\sum_{i=0}^n v_i\right) + v_{n+1}}{n+1} = \frac{\sum_{i=0}^n v_i}{n+1} + \frac{v_{n+1}}{n+1} = \\ &= \frac{\left(\sum_{i=0}^n v_i\right)n}{(n+1)n} + \frac{v_{n+1}}{n+1} = \frac{\sum_{i=0}^n v_i}{n} \cdot \frac{n}{n+1} + \frac{v_{n+1}}{n+1} = \frac{\varepsilon_n \cdot n + v_{n+1}}{n+1} \end{aligned} \quad (5)$$

1.3.2 Euler's method for differential equations

Useful later...

2 C++

2.1 Intro to `#include <random>`

2.1.1 Seed & `std::default_random_engine`

2.1.2 Distributions

2.1.2.1 `std::uniform_int_distribution<>()`

2.1.2.2 `std::uniform_real_distribution<>()`

2.1.2.3 `std::bernoulli_distribution<>()`

2.1.2.4 `std::poisson_distribution<>()`

2.1.2.5 `std::geometric_distribution<>()`

2.1.2.6 `std::discrete_distribution<>()`

2.2 Intro to data structures

2.2.1 `std::vector<T>()`

2.2.2 `std::deque<T>()`

2.2.3 Sets

2.2.4 Maps

2.3 I/O

2.3.1 `#include <iostream>`

2.3.2 Files

3 Exercises

Each exercise has 4 digits xxxx that are the same as the ones in the software folder in the course material.

3.1 First examples (1000)

Now we have to put together our **formal definitions** and our C++ knowledge to build some simple DTMCs and networks.

3.1.1 A simple Markov Chain (1100)

Let's begin our modeling journey by implementing a DTMC M s.t.

- $U = \{()\}$ it takes no input
- $X = [0, 1] \times [0, 1]$ it has infinite states (all the pairs of real numbers between 0 and 1)
- $Y = [0, 1] \times [0, 1]$
- $p : X \times X \times U \rightarrow X = \mathcal{U}(0, 1) \times \mathcal{U}(0, 1)$
- $g : X \rightarrow Y : (r_0, r_1) \mapsto (r_0, r_1)$ it outputs the current state
- $X(0) = (0, 0)$

```
#include <fstream>
#include <random>

using real_t = double;

int main() {
    std::random_device random_device;
    std::default_random_engine random_engine(random_device());
    std::uniform_real_distribution<real_t> uniform(0, 1); ①

    const size_t HORIZON = 10; ②
    std::vector<real_t> state(2, 0); ③

    for (size_t time = 0; time <= HORIZON; time++)
        for (auto &r : state)
            r = uniform(random_engine); ④

    return 0;
}
```

Listing 1: software/1100/main.cpp

3.1.2 Connect Markov Chains pt.1 (1200)

In this exercise we build 2 markov chains, and connect them...

Now let's model a system with two DTMCs M_0, M_1 , and let's define the functions

$$U(M_0, t+1) = \dots U(M_1, t+1) = \dots$$

3.1.3 Connect Markov Chains pt.2 (1300)

The same as above, but with a different connection

3.1.4 Connect Markov Chains pt.3 (1400)

The same as above, but with a different connection

3.2 Traffic light (2100, 2200, 2300)

This is

3.3 Control center (3000)

3.3.1 No network (3100)

3.3.2 Network monitor (no faults) (3200)

3.3.3 Network monitor (faults, no repair) (3300)

3.3.4 Network monitor (faults, repair) (3400)

3.3.5 Network monitor (faults, repair, correct protocol) (3500)

3.4 Statistics (4000)

3.4.1 Expected value (4100)

3.4.2 Probability (4200)

3.5 Transition matrix

3.6 Complex systems

3.6.1 Insulin pump

3.6.2 Buffer

3.6.3 Server

4 Exam

4.1 Development team (time & cost)

4.2 Backend load balancing

4.3 Heater simulation

Bibliography

- [1] Marcin Kolny, “Scaling up the Prime Video Audio-Video Monitoring Service and Reducing Costs by 90%.” Accessed: Mar. 25, 2024. [Online]. Available: <https://web.archive.org/web/20240325042615/https://www.primevideotech.com/video-streaming/scaling-up-the-prime-video-audio-video-monitoring-service-and-reducing-costs-by-90#expand>