

Software Engineering

Cicio Ionuț

23/12/2024

Contents

1 Software models	4
1.1 The “Amazon Prime Video” article	4
1.2 Formal notation	5
1.2.1 The concept of time	5
1.2.2 Markov Chain	5
1.2.3 DTMC (Discrete Time Markov Chain)	5
1.2.3.1 An example of DTMC	7
1.2.4 Network of Markov Chains	7
1.3 Tips and tricks	7
1.3.1 Mean	7
1.3.2 Eulero’s method for differential equations	8
2 C++	9
2.1 Useful <code>#include <random></code> features not in C	9
2.1.1 <code>random()</code> vs <code>random_device</code> vs <code>default_random_engine</code>	9
2.1.2 Distributions	10
2.1.2.1 <code>std::uniform_int_distribution</code>	10
2.1.2.2 <code>std::uniform_real_distribution</code>	10
2.1.2.3 <code>std::bernoulli_distribution</code>	10
2.1.2.4 <code>std::poisson_distribution</code>	10
2.1.2.5 <code>std::geometric_distribution</code>	10
2.1.2.6 <code>std::discrete_distribution</code>	10
2.2 Dynamic structures	10
2.2.1 <code>std::vector<T>()</code> instead of <code>malloc()</code>	10
2.2.2 <code>std::deque<T>()</code>	10
2.2.3 Sets	11
2.2.4 Maps	11
2.3 I/O	11
2.3.1 <code>#include <iostream></code>	11
2.3.2 Files	11
3 Exercises	12
3.1 [1000] First examples	12
3.1.1 [1100] A simple Markov Chain	12
3.1.2 [1200] Connect Markov Chains pt.1	12
3.1.3 [1300] Connect Markov Chains pt.2	13
3.1.4 [1400] Connect Markov Chains pt.3	13
3.2 [2100, 2200, 2300] Traffic light	13
3.3 [3000] Control center	13
3.3.1 [3100] No network	13
3.3.2 [3200] Network monitor (no faults)	13
3.3.3 [3300] Network monitor (faults, no repair)	13
3.3.4 [3400] Network monitor (faults, repair)	13

3.3.5 [3500] Network monitor (faults, repair, correct protocol)	13
3.4 [4000] Statistics	13
3.4.1 [4100] Expected value	13
3.4.2 [4200] Probability	13
3.5 [5000] Transition matrix	13
3.6 [6000] Complex systems	13
3.6.1 [6100] Insulin pump	13
3.6.2 [6200] Buffer	13
3.6.3 [6300] Server	13
4 Exam	13
4.1 Development team (time & cost)	14
4.2 Backend load balancing	14
4.3 Heater simulation	14
5 CASE library	14
Bibliography	15

1 Software models

When designing **complex software** we have to make major **design choices** at the beginning of the project. Often those choices can't be driven by experience or reasoning alone, that's why a **model** of the project is needed to compare different solutions. Our formal tool of choice is the **Discrete Time Markov Chain** (Section 1.2.3).

1.1 The “Amazon Prime Video” article

If you were tasked with designing the software architecture for **Amazon Prime Video** (*a live streaming service for Amazon*), how would you go about it? What if you had the **non-functional requirement** to keep the costs as low as possible?

In a recent article, Marcin Kolny, a Senior SDE at Prime Video, describes how they “*reduced the cost of the audio/video monitoring infrastructure by 90%*” [1] by using a monolith application instead of distributed microservices (an outcome one wouldn't usually expect).

While there isn't always definitive answer, one way to go about this kind choice is building a model of the system to compare the solutions. In the case of Prime Video, “*the audio/video monitoring service consists of three major components:*” [1]

- the **media converter** converts input audio/video streams
- the **defect detectors** execute algorithms that analyze frames and audio buffers in real-time looking for defects and send notifications
- the **orchestrator** controls the flow in the service

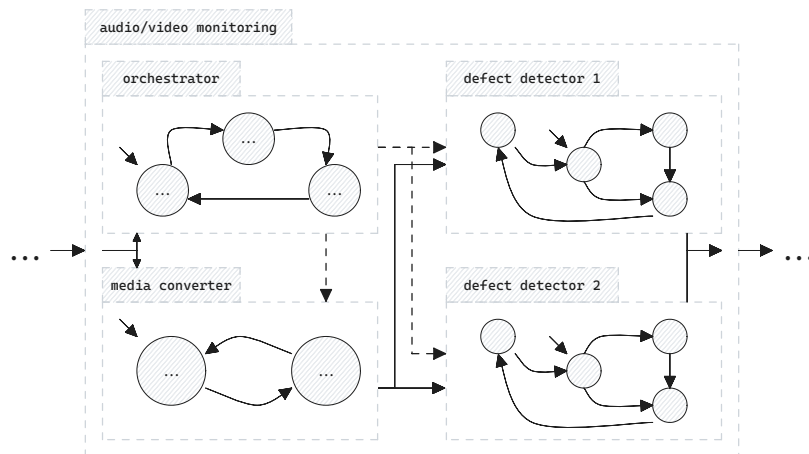


Figure 1: Model of the audio/video monitoring system

We want to model the components as some kind of stateful machine (with inputs and outputs, Figure 1), interconnect them and **simulate** the behaviour of the system as a whole.

1.2 Formal notation

1.2.1 The concept of time

TODO: rewrite

The models treated in the course evolve through **time**. Time can be modeled in many ways (I guess?), but, for the sake of simplicity, we will consider discrete time. Let W be the '*weather system*' and D the '*driving ability*' system in Section 1.2, we can define the evolution of D as

$$\begin{aligned} D(0) &= \text{'good'} \\ D(t+d) &= f(D(t), W(t)) \end{aligned} \tag{1}$$

Given a time instant t (let's suppose 12:32) and a time interval d (1 minute), the driving ability of D at 12:33 depends on the driving ability of D at the time 12:32 and the weather at 12:32.

1.2.2 Markov Chain

A Markov Chain is...

1.2.3 DTMC (Discrete Time Markov Chain)

A DTMC M is a tuple (U, X, Y, p, g) s.t.

- $U \neq \emptyset \wedge X \neq \emptyset \wedge Y \neq \emptyset$ (*otherwise stuff doesn't work*)
- U can be either
 - $\{u_1, \dots, u_n\}$ where u_i is an input value
 - $\{()\}$ if M doesn't take any input
- $X = \{x_1, \dots, x_n\}$ where x_i is a state
- $Y = \{y_1, \dots, y_n\}$ where y_i is an output value
- $p : X \times X \times U \rightarrow [0, 1]$ is the transition function
- $g : X \rightarrow Y$ is the output function

$$\forall x \in X \quad \forall u \in U \quad \sum_{x' \in X} p(x'|x, u) = 1 \tag{2}$$

$$\begin{aligned} M(0) &= x_1 \\ M(t+d) &= \begin{cases} x_1 & \text{with probability } p(x_1|M(t), U(t)) \\ x_2 & \text{with probability } p(x_2|M(t), U(t)) \\ \dots & \end{cases} \end{aligned} \tag{3}$$

It's interesting to notice that the transition function depends on the input values. If you consider the '*driving ability*' system in Section 1.2, you can see that the probability to go from **good** to **bad** is higher if the weather is rainy and lower if it's sunny.

1.2.3.1 An example of DTMC

Let's consider the development process of a team. We can define a DTMC $M = (U, X, Y, p, g)$ s.t.

- $U = \{()\}$, as it doesn't have any input
- $X = \{0, 1, 2, 3\}$
- $Y = \text{Cost} \times \text{Duration (in months)}$

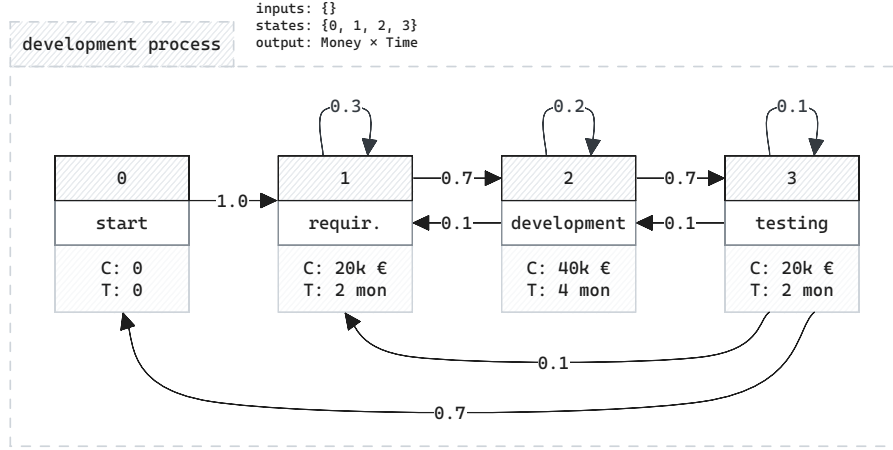


Figure 2: the model of a team's development process

$$g(x) = \begin{cases} (0, 0) & \text{if } x = 0 \\ (20000, 2) & \text{if } x = 1 \\ (40000, 4) & \text{if } x = 2 \\ (20000, 2) & \text{if } x = 3 \end{cases} \quad (4)$$

1.2.4 Network of Markov Chains

TODO...

1.3 Tips and tricks

1.3.1 Mean

TODO: 'mean' trick, ggwp

$$\begin{aligned} \varepsilon_n &= \frac{\sum_{i=0}^n v_i}{n} \\ \varepsilon_{n+1} &= \frac{\sum_{i=0}^{n+1} v_i}{n+1} = \frac{\left(\sum_{i=0}^n v_i\right) + v_{n+1}}{n+1} = \frac{\sum_{i=0}^n v_i}{n+1} + \frac{v_{n+1}}{n+1} = \quad (5) \\ \frac{\left(\sum_{i=0}^n v_i\right)n}{(n+1)n} + \frac{v_{n+1}}{n+1} &= \frac{\sum_{i=0}^n v_i}{n} \cdot \frac{n}{n+1} + \frac{v_{n+1}}{n+1} = \frac{\varepsilon_n \cdot n + v_{n+1}}{n+1} \end{aligned}$$

1.3.2 Euler's method for differential equations

Useful later...

2 C++

This section will cover the basics needed for the exam.

2.1 Useful `#include <random>` features not in C

The `<random>` library offers useful tools to build our models. It makes Markov Chains and probabilistic models really easy to implement.

2.1.1 `random()` vs `random_device` vs `default_random_engine`

In C++ there are many ways to **generate random numbers**, I'm gonna keep it short and sweet: don't use `random()`, use `std::random_device` to generate the **seed** and from then on just some random engine like `std::default_random_engine`.

```
#include <iostream>
#include <random>

int main() {
    std::cout << random() ① << std::endl;

    std::random_device random_device; ②
    std::cout << random_device() ③ << std::endl;

    std::default_random_engine r_engine(random_device() ④ );
    std::cout << r_engine() ⑤ << std::endl;
}
```

Listing 1: examples/random.cpp

I'll explain: `random()` ① doesn't work very well (TODO: link to the article*), you're encouraged to use the generators C++ offers... but each works in a different way, and you have to choose the best one depending on your needs. `std::random_device` ② is used to generate the **seed** ④ because it uses device randomness (if available, TODO: link to docs / article *), but it's really slow. That **seed** is used to to instantiate one of the other engines ⑤, like `::default_random_engine` (TODO: link with different engines and types). TODO BONUS: only `r_eng` used with distr

Note:

If you see ③ and ⑤, `random_device` and `r_engine` aren't functions, they are instances, but you can call the `()` operator on them because C++ has operator overloading, and it allows you to call custom operators on instances... the same goes for `std::cout` and `<<`

2.1.2 Distributions

Just the capability of generating random numbers isn't enough, we often need to manipulate those numbers to fit our needs. Luckily, C++ covers basically all of them... for example, this is how easy it is to simulate a transition matrix in Figure 2:

```
-----
#include <iostream>
#include <random>

const size_t HORIZON = 15;

int main() {
    std::random_device random_device;
    std::default_random_engine random_engine(random_device());

    std::discrete_distribution<size_t> transition_matrix[] = {
        {0, 1},
        {0, .3, .7},
        {0, .2, .2, .6},
        {0, .1, .2, .1, .6},
        {1},
    };

    size_t state = 0;
    for (size_t time = 0; time < HORIZON; time++) {
        state = transition_matrix[state](random_engine);
        std::cout << state << std::endl;
    }
}
-----
```

Listing 2: examples/transition_matrix.cpp

2.1.2.1 std::uniform_int_distribution

2.1.2.2 std::uniform_real_distribution

2.1.2.3 std::bernoulli_distribution

2.1.2.4 std::poisson_distribution

2.1.2.5 std::geometric_distribution

2.1.2.6 std::discrete_distribution

2.2 Dynamic structures

2.2.1 std::vector<T>() instead of malloc()

2.2.2 std::deque<T>()

2.2.3 Sets

2.2.4 Maps

2.3 I/O

2.3.1 `#include <iostream>`

2.3.2 Files

3 Exercises

Each exercise has 4 digits xxxx that are the same as the ones in the software folder in the course material.

3.1 [1000] First examples

Now we have to put together our **formal definitions** and our C++ knowledge to build some simple DTMCs and networks.

3.1.1 [1100] A simple Markov Chain

Let's begin our modeling journey by implementing a DTMC M s.t.

- $U = \{()\}$ it takes no input
- $X = [0, 1] \times [0, 1]$ it has infinite states (all the pairs of real numbers between 0 and 1)
- $Y = [0, 1] \times [0, 1]$
- $p : X \times X \times U \rightarrow X = \mathcal{U}(0, 1) \times \mathcal{U}(0, 1)$
- $g : X \rightarrow Y : (r_0, r_1) \mapsto (r_0, r_1)$ it outputs the current state
- $X(0) = (0, 0)$

```
#include <fstream>
#include <random>

using real_t = double;

int main() {
    std::random_device random_device;
    std::default_random_engine random_engine(random_device());
    std::uniform_real_distribution<real_t> uniform(0, 1); ①

    const size_t HORIZON = 10; ②
    std::vector<real_t> state(2, 0); ③

    for (size_t time = 0; time <= HORIZON; time++)
        for (auto &r : state)
            r = uniform(random_engine); ④

    return 0;
}
```

Listing 3: software/1100/main.cpp

3.1.2 [1200] Connect Markov Chains pt.1

In this exercise we build 2 markov chains, and connect them...

Now let's model a system with two DTMCs M_0, M_1 , and let's define the functions

$$U(M_0, t+1) = \dots U(M_1, t+1) = \dots$$

3.1.3 [1300] Connect Markov Chains pt.2

The same as above, but with a different connection

3.1.4 [1400] Connect Markov Chains pt.3

The same as above, but with a different connection

3.2 [2100, 2200, 2300] Traffic light

This is

3.3 [3000] Control center

3.3.1 [3100] No network

3.3.2 [3200] Network monitor (no faults)

3.3.3 [3300] Network monitor (faults, no repair)

3.3.4 [3400] Network monitor (faults, repair)

3.3.5 [3500] Network monitor (faults, repair, correct protocol)

3.4 [4000] Statistics

3.4.1 [4100] Expected value

3.4.2 [4200] Probability

3.5 [5000] Transition matrix

3.6 [6000] Complex systems

3.6.1 [6100] Insulin pump

3.6.2 [6200] Buffer

3.6.3 [6300] Server

4 Exam

4.1 Development team (time & cost)

4.2 Backend load balancing

4.3 Heater simulation

5 CASE library

TODO...

Bibliography

- [1] Marcin Kolny, “Scaling up the Prime Video Audio-Video Monitoring Service and Reducing Costs by 90%.” Accessed: Mar. 25, 2024. [Online]. Available: <https://web.archive.org/web/20240325042615/https://www.primevideotech.com/video-streaming/scaling-up-the-prime-video-audio-video-monitoring-service-and-reducing-costs-by-90#expand>